

# Parallel Image Processing with CUDA

## – A case study with the Canny Edge Detection Filter –

Daniel Weingaertner

Informatics Department  
Federal University of Paraná - Brazil

Hochschule Regensburg  
02.05.2011

# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)
- 3 GPGPU and CUDA
- 4 Integrating CUDA and ITK
- 5 Canny Edge Detection
- 6 Experimental Results
- 7 Conclusion

# Paraná – Brazil




## Brazil – Europe



# Paraná



# Curitiba



# Federal University of Paraná



## Informatic Department



# Informatics Department

Undergraduate: Bachelor in Computer Science

- 8 semesters course
- 80 incoming students per year

Bachelor in Biomedical Informatics

- 8 semesters course
- 30 incoming students per year

Graduate: Master and PhD in Computer Science

- Algorithms, Image Processing, Computer Vision, Artificial Intelligence
- Databases, Scientific Computing and Open Source Software, Computer-Human Interface
- Computer Networks, Embedded Systems




# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)**
- 3 GPGPU and CUDA
- 4 Integrating CUDA and ITK
- 5 Canny Edge Detection
- 6 Experimental Results
- 7 Conclusion

# Insight Toolkit (ITK)

- Created in 1999, Open Source, Multi platform, Object Oriented (Templates), Good documentation and support



**Figure:** Image Processing Workflow in ITK

# ITK - Sample code

---

```
1 #include "itkImage.h"
2 #include "itkImageFileReader.h"
3 #include "itkImageFileWriter.h"
4 #include "itkCannyEdgeDetectionImageFilter.h"
5
6 typedef itk::Image<float,2> ImageType;
7 typedef itk::ImageFileReader< ImageType > ReaderType;
8 typedef itk::ImageFileWriter< ImageType > WriterType;
9 typedef itk::CannyEdgeDetectionImageFilter< ImageType, ImageType > CannyFilter;
10
11 int main (int argc, char** argv){
12
13     ReaderType::Pointer reader = ReaderType::New();
14     reader->SetFileName( argv[1] );
15     reader->Update();
16
17     CannyFilter::Pointer canny = CannyFilter::New();
18     canny->SetInput( reader->GetOutput() );
19     canny->SetVariance( atof( argv[3] ) );
20     canny->SetUpperThreshold( atoi( argv[4] ) );
21     canny->SetLowerThreshold( atoi( argv[5] ) );
22     canny->Update();
23
24     WriterType::Pointer writer = WriterType::New();
25     writer->SetFileName( argv[2] );
26     writer->SetInput( canny->GetOutput() );
27     writer->Update();
28
29     return EXIT_SUCCESS;
30 }
```


---

# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)
- 3 GPGPU and CUDA**
- 4 Integrating CUDA and ITK
- 5 Canny Edge Detection
- 6 Experimental Results
- 7 Conclusion


# What is GPGPU Computing?

- The use of the GPU for general purpose computation
- CPU and GPU can be used concurrently
- To the end user, its simply a way to run applications faster.

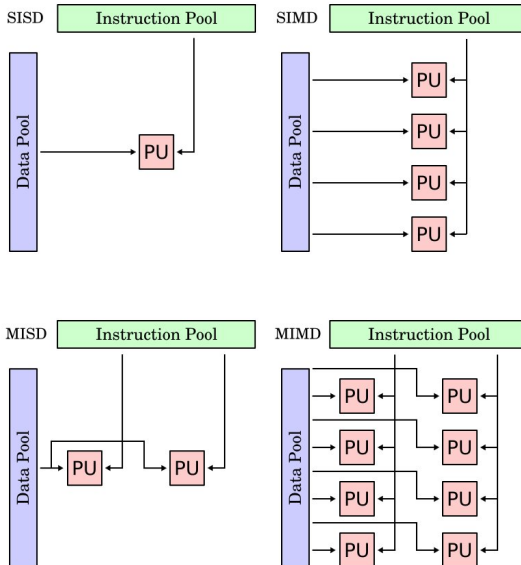


# What is CUDA?

- CUDA = *Compute Unified Device Architecture*.
- *General-Purpose Parallel Computing Architecture*.
- Provides libraries, C language extension and hardware driver.




# Parallel Processing Models



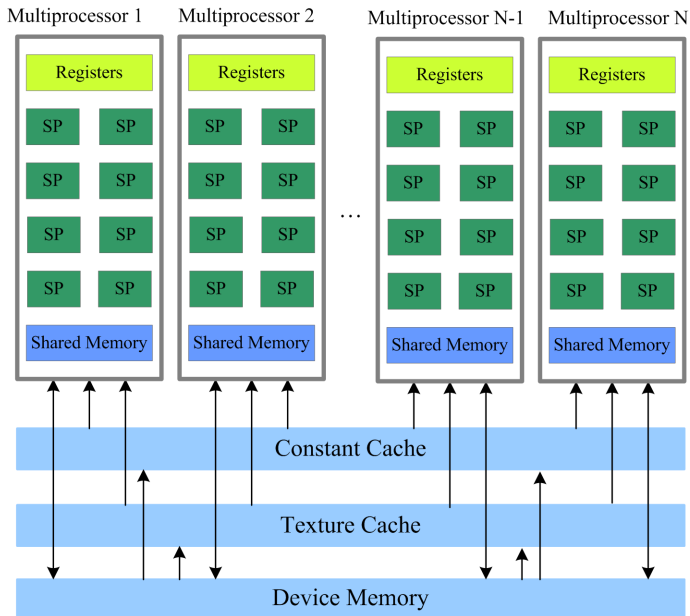
# Single-Instruction Multiple-Thread Unit

- Creates, handles, schedules and executes groups of 32 threads (*warp*).
- All threads in a *warp* start at the same point.
- But they are “free” to jump to different code positions independently.





# CUDA Architecture Overview



# Optimization Strategies for CUDA

## Main optimization strategies for CUDA involve:

- Optimized/careful memory access
- Maximization of processor utilization
- Maximization of non-serialized instructions

# CUDA - Sample Code

---

```
1 #include <stdio.h>
2 #include <assert.h>
3 #include <cuda.h>
4 void incrementArrayOnHost(float *a, int N)
5 {
6     int i;
7     for (i=0; i < N; i++) a[i] = a[i]+1.f;
8 }
9 __global__ void incrementArrayOnDevice(float *a, int N)
10 {
11     int idx = blockIdx.x*blockDim.x + threadIdx.x;
12     if (idx<N) a[idx] = a[idx]+1.f;
13 }
14 int main(void)
15 {
16     float *a_h, *b_h;           // pointers to host memory
17     float *a_d;                 // pointer to device memory
18     int i, N = 10000;
19     size_t size = N*sizeof(float);
20     a_h = (float *)malloc(size);
21     b_h = (float *)malloc(size);
22     cudaMalloc((void **) &a_d, size);
23     for (i=0; i<N; i++) a_h[i] = (float)i;
24     cudaMemcpy(a_d, a_h, sizeof(float)*N, cudaMemcpyHostToDevice);
25     incrementArrayOnHost(a_h, N);
26     int blockSize = 256;
27     int nBlocks = N/blockSize + (N%blockSize == 0?0:1);
28     incrementArrayOnDevice <<< nBlocks, blockSize >>> (a_d, N);
29     cudaMemcpy(b_h, a_d, sizeof(float)*N, cudaMemcpyDeviceToHost);
30     free(a_h); free(b_h); cudaFree(a_d);
31 }
```

---

# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)
- 3 GPGPU and CUDA
- 4 Integrating CUDA and ITK**
- 5 Canny Edge Detection
- 6 Experimental Results
- 7 Conclusion

# Integrating CUDA Filters into ITK Workflow

## ITK community suggests:

- Re-implement filters where parallelizing provides significant speedup
- Consider the entire workflow: copying to/from the GPU is very time consuming

## Careful!

“Premature optimization is the root of all evil!” (Donald Knuth)

# Integrating CUDA Filters into ITK Workflow

## ITK community suggests:

- Re-implement filters where parallelizing provides significant speedup
- Consider the entire workflow: copying to/from the GPU is very time consuming


## Careful!

“Premature optimization is the root of all evil!” (Donald Knuth)

# CUDA Insight Toolkit (CITK)

## Changes to ITK

- Slight architecture change: CudaImportImageContainer
- Backwards compatible
- Data transfer between HOST and DEVICE only “on demand”
- Allows for filter chaining inside the DEVICE



# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)
- 3 GPGPU and CUDA
- 4 Integrating CUDA and ITK
- 5 Canny Edge Detection**
- 6 Experimental Results
- 7 Conclusion



- itkCudaCannyEdgeDetectionImageFilter

---

**Algorithm 1** Canny Edge Detection Filter

---

Gaussian Smoothing

Gradient Computation

Non-Maximum Supression

Histeresis

---

# Gradient Computation with Sobel Filter

- itkCudaSobelEdgeDetectionImageFilter

<b>-1</b>	<b>0</b>	<b>1</b>
<b>-2</b>	<b>0</b>	<b>2</b>
<b>-1</b>	<b>0</b>	<b>1</b>

(a) Sobel X

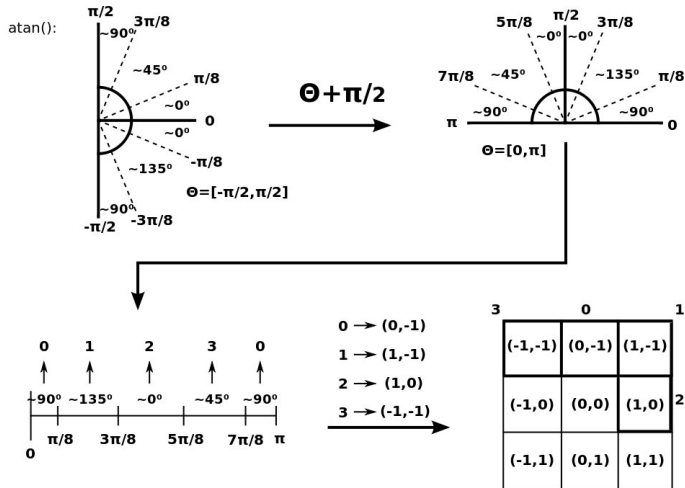
<b>1</b>	<b>2</b>	<b>1</b>
<b>0</b>	<b>0</b>	<b>0</b>
<b>-1</b>	<b>-2</b>	<b>-1</b>

(b) Sobel Y

$$L_v = \sqrt{L_x^2 + L_y^2} \quad (1)$$

$$\theta = \arctan \left( \frac{L_y}{L_x} \right) \quad (2)$$

# Optimization for Edge Direction Computation



# Code Extract from CudaSobel

```
if ((pos.x) && ((size.x-1)-pos.x) && (pos.y) && ((size.y-1)-pos.y)){

    diagonal.x = tex1Dfetch(texRef, (pixIdx-size.x-1));
    diagonal.y = tex1Dfetch(texRef, (pixIdx-size.x+1));
    diagonal.z = tex1Dfetch(texRef, (pixIdx+size.x-1));
    diagonal.w = tex1Dfetch(texRef, (pixIdx+size.x+1));
    cross.x = tex1Dfetch(texRef, (pixIdx-size.x));
    cross.y = tex1Dfetch(texRef, (pixIdx+size.x));
    cross.z = tex1Dfetch(texRef, (pixIdx-1));
    cross.w = tex1Dfetch(texRef, (pixIdx+1));

    /// SobelX
    g_i.x -= (diagonal.x+cross.z+cross.z+diagonal.z);
    g_i.x += (diagonal.y+cross.w+cross.w+diagonal.w);

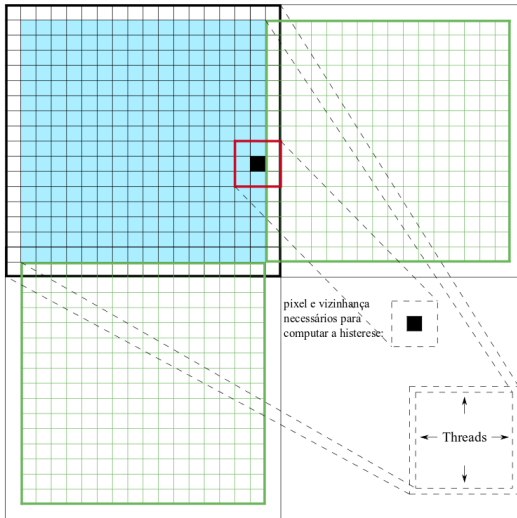
    /// SobelY
    g_i.y -= (diagonal.z+cross.y+cross.y+diagonal.w);
    g_i.y += (diagonal.x+cross.x+cross.x+diagonal.y);
}

Magnitude[pixIdx] = sqrtf((g_i.x*g_i.x) + (g_i.y*g_i.y));

theta = (g_i.x != 0)*(int)(atanf(__fdivdef(g_i.y,g_i.x)*__fdivdef(180,M_PI)) + 90;
if (theta > 157) theta -= 158;
theta = ceilf(__fdivdef(theta-22,45));
Direction[pixIdx] = make_short2(1-(theta == 0)-((theta == 1)<<1),(theta == 2)-1);

}
```

# Hysteresis Operation



# Hysteresis Algorithm

---

## Algorithm 2 Hysteresis on CPU

---

Transfers the Gradient/NMS images to the GPU

**repeat**

    Run the hysteresis kernel on GPU

**until** no pixel changes status

Return edge image

---

# Hysteresis Algorithm

---

## Algorithm 3 Hysteresis on GPU

---

Load an image region with size 18x18 into shared memory

modified  $\leftarrow$  false

**repeat**

    modified\_region  $\leftarrow$  false

    Synchronize *threads* of same multiprocessor

**if** Pixel changes status **then**

        modified  $\leftarrow$  true

        modified\_region  $\leftarrow$  true

**end if**

    Synchronize *threads* of same multiprocessor

**until** modified\_region = false

**if** modified = true **then**

    Update modified status on HOST

**end if**

---

# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)
- 3 GPGPU and CUDA
- 4 Integrating CUDA and ITK
- 5 Canny Edge Detection
- 6 Experimental Results**
- 7 Conclusion



## Hardware:

- Server:

- ▶ CPU: 4x AMD Opteron(tm) Processor 6136 2,4GHz with 8 cores, each with 512 KB cache and 126GB RAM
- ▶ GPU1: NVidia Tesla C2050 with 448 1,15GHz cores and 3GB RAM.
- ▶ GPU2: NVidia Tesla C1060 com 240 1,3GHz cores and 4GB RAM.


- Desktop:

- ▶ CPU: Intel®Core(TM)2 Duo E7400 2,80GHz with 3072 KB cache and 2GB RAM
- ▶ GPU: NVidia GeForce 8800 GT with 112 1,5GHz cores and 512MB RAM.


- Images from the Berkeley Segmentation Dataset

Base	Image resolution	Num. of Images
B1	$321 \times 481$ e $481 \times 321$	100
B2	$642 \times 962$ e $962 \times 642$	100
B3	$1284 \times 1924$ e $1924 \times 1284$	100
B4	$2568 \times 3848$ e $3848 \times 2568$	100


# Performance Tests




# Performance Tests



# Performance Tests



# Performance Tests



# Summary

- 1 Introduction
- 2 Insight Toolkit (ITK)
- 3 GPGPU and CUDA
- 4 Integrating CUDA and ITK
- 5 Canny Edge Detection
- 6 Experimental Results
- 7 Conclusion**

# Conclusion

## Parallel Programming

- Parallel programming is definitely the way to go.
- Implement efficient parallel code is demanding.
- Programmer should know more details about the hardware, especially memory architecture.

## Canny Filter with CUDA

- We had a great speedup on the edge detection filter
- Also noticed that the existing implementation is not efficient
- There is still a LOT of work if we want to parallelize ITK.



# Conclusion

## Parallel Programming

- Parallel programming is definitely the way to go.
- Implement efficient parallel code is demanding.
- Programmer should know more details about the hardware, especially memory architecture.

## Canny Filter with CUDA

- We had a great speedup on the edge detection filter
- Also noticed that the existing implementation is not efficient
- There is still a LOT of work if we want to parallelize ITK.

**Thank You!**

Daniel Weingaertner  
danielw@inf.ufpr.br